

# Neural Networks 2017



Dr. Wojtek Kowalczyk

[wojtek@liacs.nl](mailto:wojtek@liacs.nl)

# Agenda



- Course objectives, organization, grading
- Neural Networks: motivation, history, context
- Deep Learning Revolution
- Statistical Pattern Recognition
- Assignment 0

# Key Objectives

---

1. Learn key concepts behind the “traditional” neural networks (***Perceptron, MLP, Backpropagation***)
2. Learn basics of ***Deep Learning and its applications***
3. Learn key techniques for ***parameter optimization*** (finding optima of functions of many variables)
4. Master modern tools for training deep networks on GPUs (***Python + Theano/Tensorflow/Keras***)
5. Gain ***practical experience*** with applying Neural Networks to various problems

# Course Organization

---

## Lectures:

theory + example applications + practical tips

Wednesdays, 11:15-13:00, rooms: 3xB02+10x312

## Computer Lab:

Wednesdays, 9:00-10:45, Rooms 302-304, 306-308

*Presence during lectures & comp. lab is **not compulsory***

# Grading

---

- Course value: 6 ects (workload: 2days/week)
- Final grade is based on two components:
  - written exam (40%):  
various questions/problems related to theory and practice.
  - practical (60%):  
three programming assignments:  
  
Assignment 0: elementary Python skills (not graded)  
Assignment 1: Linear Models + MLP  
Assignment 2: Deep Neural Networks on GPU's  
Assignment 3: ***A DL Challenge***
  - **Both parts (exam+prac) must be “passed” (grade > 5.5)!**

# Textbook

---

[www.deeplearningbook.org](http://www.deeplearningbook.org)

Available on-line; print it to PDF  
or check an example pdf dump  
(*a link will follow soon*)

The book is about 800 pages long;  
don't despair – we will use just 30%!


**This week: Chapter 1**

**Next week: Chapter 5**

(*Chapters 2-4: “background”*)

Neural Networks

NN 1



# Additional Resources

---

## Deep Learning in Matlab:

- [github.com/rasmusbergpalm/DeepLearnToolbox](https://github.com/rasmusbergpalm/DeepLearnToolbox)  
(a very concise intro to Deep Learning + Matlab code)
- [www.aston.ac.uk/eas/research/groups/ncrg/resources/netlab/](http://www.aston.ac.uk/eas/research/groups/ncrg/resources/netlab/)

## Deep Learning:

- Kevin Duh: Neural Networks and Deep Learning  
[cl.naist.jp/~kevinduh/a/deep2014/](http://cl.naist.jp/~kevinduh/a/deep2014/)
- <http://deeplearning.net/tutorial/>
- Andrew Ng on Deep Learning:  
[www.youtube.com/watch?v=n1ViNeWhC24](https://www.youtube.com/watch?v=n1ViNeWhC24)

# Neural Networks:

motivation

history

context



# How could we program computers to do:

---

- Optical Character Recognition
- Text2Speech translation
- Speech2Text translation
- Robot arm control
- Face recognition
- Driving a car
- Detecting fraud with credit cards
- Predicting \$/€ exchange rate

How people do it?



“training data” + “self-learning”

# The biological model

Human brain

around  $10^{11}$  neurons, and  $10^{15}$  interconnections (synapses)

## Structure of a Typical Neuron




# The biological model

## How a neuron works ?

- It receives input signals through its dendrites
- The input signals generate difference of electrical potential on the cell membrane
- The difference is propagated to the axon hillock
- A train of electrical impulses is generated along the axon
- The impulses on the axon generate the release in the synaptic space of some neurotransmitters

## Structure of a Typical Neuron



***Learning by changing  
the topology & thickness  
of connections***

# Basic principle

Artificial neural network (ANN)= set of highly interconnected simple processing units (also called neurons)

Processing unit = simplified model of a neuron

Learning = finding the architecture of connections and their weights

# Computer versus Brain

## Von Neumann machine

- One or a few high speed (ns) processors with considerable computing power
- One or a few shared high speed buses for communication
- Sequential memory access by address
- Problem-solving knowledge is separated from the computing component
- Hard to be adaptive

Neural Networks

## Human Brain

- Large number ( $10^{11}$ ) of low speed processors (ms) with limited computing power
- Large number ( $10^{15}$ ) of low speed connections
- Content addressable recall (CAM)
- Problem-solving knowledge resides in the connectivity of neurons
- ***Adaptation by changing the connectivity (topology & thickness of connections)***

NN 1

13

# Human Brain

---

- **graceful degradation:**  
its performance degrades gracefully under partial damage.
- **self-learning:**  
it can learn (reorganize itself) from experience.
- **self-healing:**  
partial recovery from damage is possible
- **massive parallelism:**  
performs parallel computations extremely efficiently. E.g., complex visual perception takes less than 100 ms, that is, 10 processing steps!
- **intelligence and self-awareness** (whatever it means)

# ANN History: Roots ...

Roots of Neural Networks are in:

**Neurobiological studies** (more than one century ago):

- How do nerves behave when stimulated by different magnitudes of electric current? Is there a minimal threshold needed for nerves to be activated? Given that no single nerve cell is long enough, how do different nerve cells communicate among each other?

**Psychological studies:**

- How do animals learn, forget, recognize and perform other types of tasks?

**Psycho-physical experiments**

- Understand how individual neurons and groups of neurons work.

# ANN History: First steps

## ❑ **Pitts & McCulloch (1943)**

- ❑ First mathematical model of biological neurons
- ❑ All Boolean operations can be implemented by these neuron-like nodes (with different threshold and excitatory/inhibitory connections).
- ❑ Competitor to Von Neumann model for general purpose computing device
- ❑ Origin of automata theory.

## ❑ **Hebb (1949)**

- ❑ Hebbian rule of learning: increase the connection strength between neurons  $i$  and  $j$  whenever both  $i$  and  $j$  are activated.
- ❑ Or increase the connection strength between nodes  $i$  and  $j$  whenever both nodes are simultaneously ON or OFF.



# ANN History: Perceptron

## ❑ Early booming (50's – early 60's)

- Rosenblatt (1958)
  - Perceptron: network of threshold nodes for pattern classification. Perceptron learning rule – first learning algorithm
  - Perceptron convergence theorem: everything that can be represented by a perceptron can be learned
- Widrow and Hoff (1960, 1962)
  - Learning rule that is based on minimization methods
- Minsky's attempt to build a general purpose machine with Pitts/McCulloch units

# ANN History: Setback ...

- **The setback (mid 60's – late 70's)**
  - Minsky and Papert publish a book "Perceptrons" (1969):
    - Single layer perceptrons **cannot represent** (learn) simple functions such as **XOR**
    - Multi-layer of non-linear units may have greater power but there was no learning algorithm for such nets
    - Scaling problem: connection weights may grow infinitely
  - ***US Defense/Government stop funding research on ANN***


# ANN History: Backpropagation

- **Renewed enthusiasm and progress (80's – 90's)**
  - New techniques
    - **Backpropagation** learning for multi-layer feed forward nets (with non-linear, differentiable node functions)
    - Physics inspired models (Hopfield net, Boltzmann machine, etc.)
    - Unsupervised learning (LVQ nets, Kohonen nets)
  - Impressive applications (character recognition, speech recognition, text-to-speech transformation, process control, associative memory, etc.)

## **But:**

- Criticism from Statisticians, Neurologists, Biologists, Ordinary Users, ...
- Lots of ad-hoc solutions, "wild creativity"
- A lot of rubbish is produced ...

# Supervised Learning



- **Training data:** inputs & targets (= desired outputs)
- **Network:** a complicated function with many parameters to tune
- **Training:** a process of tweaking the parameters to minimize the errors of predictions (outputs should be close to targets)
- **Generalization:** we want the network to perform well on data that was not used during the training process!

$$Error = \sum_{Examples} (output_i - target_i)^2$$

# Example: Recognizing Handwritten Digits




---

Training set: 2000 images of digits

Each image: 16x16 pixels = a vector of 256 components

The network: 256 inputs and 10 outputs

Test set: 1000 images of digits



# Backpropagation algorithm

---

- The backpropagation algorithm searches for weight values that **minimize the total error of the network**
- It consists of the **repeated application of two passes**:
  - **Forward pass**: in this step the network is activated on one example and the **error** of each neuron of the output layer is **computed**.
  - **Backward pass**: in this step the network error is used for updating the weights. Starting at the output layer, **the error is propagated backwards through the network, layer by layer**, with help of the **generalized delta rule**. Finally, all weights are updated.
- **No guarantees of convergence**  
(when learning rate too big or too small)
- In case of convergence: **local (or global) minimum**
- In practice: **try several starting configurations and learning rates**.

# Three examples of MLP



**NetTalk:** a network that reads aloud texts

**ALVINN:** a Neural network that drives a car

**Falcon:** a real-time system for detecting fraud  
with credit card transactions


# NETtalk (Sejnowski & Rosenberg, 1987)

---

- The task is to **learn to pronounce English text** from examples (text-to-speech)
- Training data: a list of **<phrase, phonetic representation>**
- **Input:** 7 consecutive characters from written text presented in a moving window that scans text
- **Output:** phoneme code giving the pronunciation of the letter at the center of the input window
- **Network topology:** 7x29 binary inputs (26 chars + punctuation marks), 80 hidden units and 26 output units (phoneme code). Sigmoid units in hidden and output layer



# NETtalk




## NETtalk (contd.)

---

- Training protocol: 95% accuracy on training set after 50 epochs of training by full gradient descent.  
78% accuracy on a test set
- DEC-talk: a rule-based system crafted by experts (a decade of efforts by many linguists)
- Functionality/Accuracy almost the same
- Try: <http://cnl.salk.edu/Media/nettalk.mp3>

# ALVINN (1989)

D.A. Pomerleau, Autonomous Land Vehicle In a Neural Network (NIPS '89)



# FALCON: A Fraud Detection System

---

<http://www.fico.com/en/Products/DMApps/Pages/FICO-Falcon-Fraud-Manager.aspx>



Right now, leading institutions around the world trust Fair Isaac's Falcon™ Fraud Manager to protect more than **450 million active credit and debit cards**. Why? They know they'll receive regular technological advances which will allow them to stay a step ahead of new and emerging fraud types. **Protecting 65% of the world's credit card transactions, Falcon detects fraud** with pinpoint accuracy via **proven neural network models** and other proprietary predictive technologies. Debit, credit, oil and retail card issuers in numerous marketplaces rely on Falcon to detect and stop fraudulent transactions - and combat identity theft - in real time.

# Why only “shallow networks” ?

---

- In practice, "classical" multi-layer networks work fine only for a very small number of hidden layers (typically 1 or 2) - this is an empirical fact ...
- Adding layers is harmful because:
  - the increased number of weights quickly leads to data overfitting (lack of generalization)
  - huge number of (bad) local minima trap the gradient descent algorithm
  - vanishing or exploding gradients (the update rule involves products of many numbers) cause additional problems

**Overfitting:**  
which polynomial  
better approximates  
the green line?



The more parameters  
the higher the risk  
of overfitting !

# Why do we need many layers?




- In theory, one hidden layer is sufficient to model any function with an arbitrary accuracy; however, the number of required nodes and weights grows exponentially fast
- The deeper the network the less nodes are required to model "complicated" functions
- Consecutive layers "learn" features of the training patterns; from simplest (lower layers) to more complicated (top layers)
- Visual cortex consists of about 10 layers

# An end of hype?

---

## Dominance of new techniques (90's - 2005):

- Support Vector Machines (Vapnik)
- Kernel Methods (Vapnik, Scholkopf, ...)
- Ensemble methods (Breiman, Hasti, Tibshirani, Friedman,...)
- Random Forests
- *Neural Networks lose their prominent role in the fields of Pattern Recognition and Machine Learning*
- *since 1996 no new textbooks on Neural Networks !*





# Deep Learning Revolution



**1989 : LeCun et al: A Convolutional Network for OCR (AT&T)**

**~ 2006 : G. Hinton's group: new ideas, inventions, applications**

- Restricted Boltzmann Machine as building block for DNN
- Contrastive Divergence algorithm
- Combine supervised with unsupervised learning
- Deep Belief Networks
- Stacked Auto-Encoders
- Deep Recurrent Networks

**Many spectacular applications beating existing approaches**

# Enabling Factors



- **Availability of “Big Data”:** the more data we have the better we can train a network
- **Powerful Hardware (GPU’s):** speedup of the training process by 100-1000 times reduces the training time from years to hours
- **New algorithms and architectures:** leaving the MLP standard behind ...
- ***Many spectacular successes!***

# Key Architectures

---

- **Convolutional Networks:** when adding layers enforce hidden nodes to learn "local features" - that reduces the number of parameters
- **Autoencoders:** add layers one by one, training them separately to learn a hierarchy of features
- **Deep Recurrent Networks:** networks for modeling sequential data
- **Restricted Boltzmann Machines:** a generative model (a pdf) of the training data

# Autonomous Land Vehicle


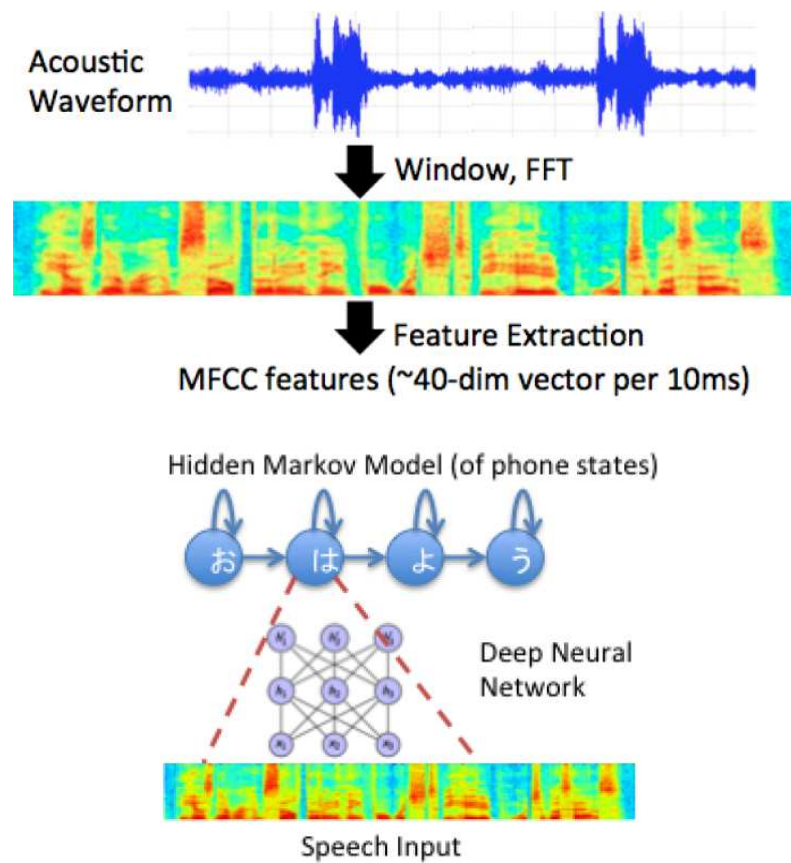
(DARPA's GrandChallenge 2005 contest)





# Other success stories

<http://cl.naist.jp/~kevinduh/a/deep2014/140121-ResearchSeminar.pdf>



# LeNet: OCR system for digits

---



training:  
60.000 images

testing:  
10.000 images

each image:  
32x32 pixels


accuracy: 99.7%

# A Convolutional Filter

Let us suppose that in an input image we want to find locations that look like a 3x3 cross. Define a matrix  $F$  (called a **receptive field** or **convolutional filter**) and "convolve it" with all possible locations in the image. We will get another (smaller) matrix with "degrees of overlap":




# LeNet5



- Input: 32x32 pixel image
- Cx: Convolutional layer
- Sx: Subsample layer  
(reduces image size by averaging 2x2 patches)
- Fx: Fully connected layer

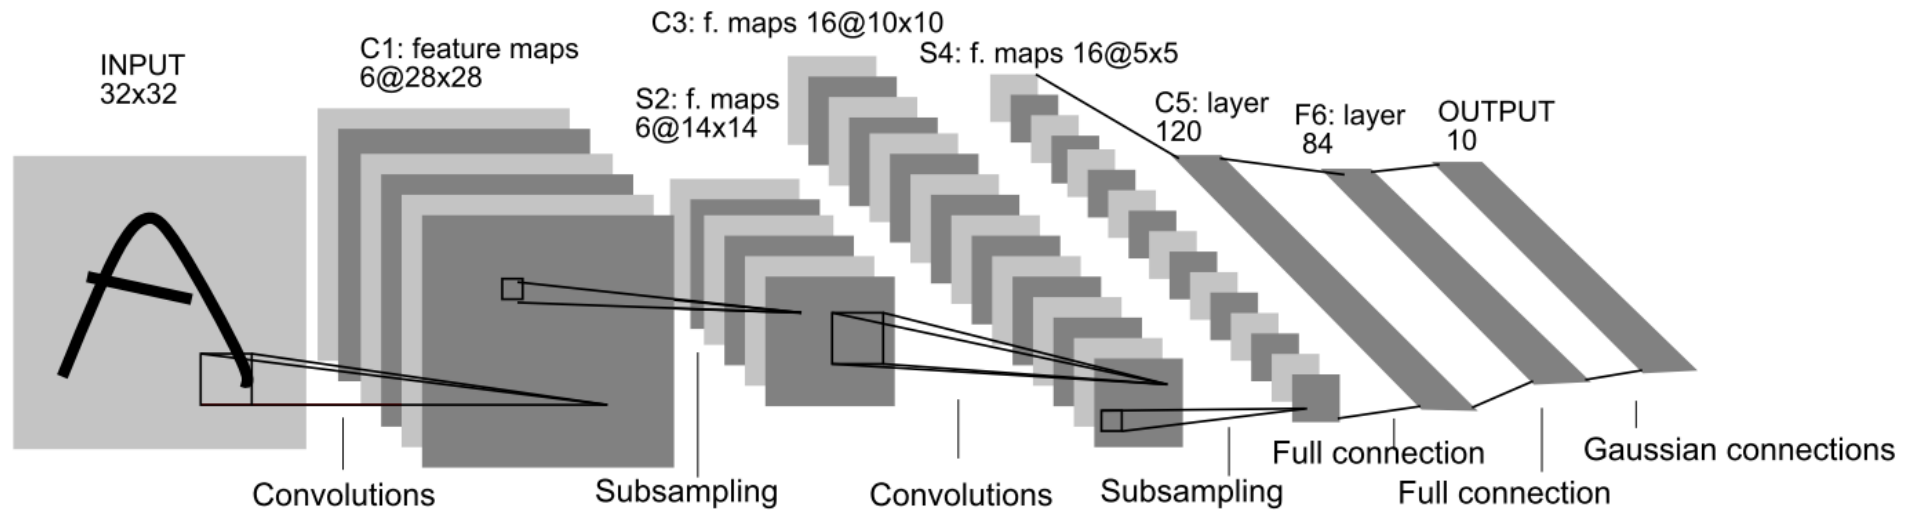


# LeNet 5: Layer C1




- C1: Convolutional layer with 6 feature maps of size 28x28.
- Each unit of C1 has a 5x5 receptive field in the input layer.
- Shared weights  $(5*5+1)*6=$ **156 parameters to learn**  
Connections:  $28*28*(5*5+1)*6=122304$
- If it was fully connected we had:  
 $(32*32+1)*(28*28)*6 =$ **4.821.600 parameters**

# LeNet 5: Layer S2



- S2: Subsampling layer with 6 feature maps of size 14x14 2x2 nonoverlapping receptive fields in C1
- Layer S2:  $6 \times 2 = 12$  trainable parameters.
- Connections:  $14 \times 14 \times (2 \times 2 + 1) \times 6 = 5880$

# LeNet 5: totals




- The whole network has:
  - 1256 nodes
  - 64.660 connections
  - **9.760 trainable parameters (and not millions!)**

# Restricted Boltzmann Machine

---

*Hidden  
layer ( $h$ )*



*Visible  
layer ( $x$ )*

- 2-layered network
- $W=(w_{ij})$  : matrix of weights  
 $b = (b_i)$  : biases of visible layer  
 $c = (c_j)$  : biases of hidden layer
- nodes take values 0 or 1
- **model of probability distr. of  $P(x, h)$**
- non-deterministic behavior:

$$P(h_i = 1|x) = \text{sigm}(c_i + W_i x)$$

$$P(x_i = 1|h) = \text{sigm}(b_i + W_i h)$$

# RBM as a model of $P(x, h)$

---

$P(x, y)$  is defined as:

$$P(x, h) = \frac{e^{-E(x, h)}}{Z}$$

where **energy**  $E$  and **normalization**  $Z$  are given by:

$$E(x, h) = -b'x - c'h - h'Wx$$

$$Z(x, h) = \sum_{x, h} e^{-E(x, h)}$$

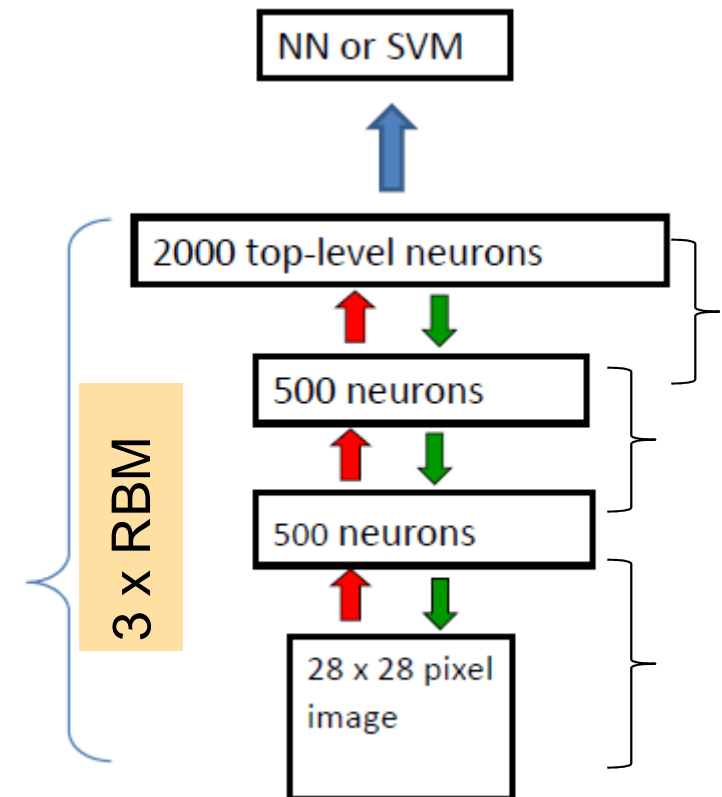
# Key properties of RBMs

---

- "unsupervised learning" (no target labels needed)
- Hidden layer learns features that contain information needed to reconstruct the input vectors (feature extractor)
- RBMs can be stacked to form deep networks, with layers trained one after another; they learn more and more complex features (hierarchy of features)
- RBMs are building blocks for deep belief networks, and deep autoencoders

# Deep Belief Networks

- the first 3 layers are RBMs
- they are trained one after another with Contrastive Divergence (unsupervised)
- next, an additional output layer is added and the whole network is trained with backpropagation
- Alternatively, instead of adding the final layer one can add a shallow MLP or a conventional classifier (e.g., SVM) and train it with conventional (supervised) methods



# More...




- Why DBN is better than MLP?
  - in MLP “delta’s” vanish (or explode) from layer to layer
  - RBMs “discover” intrinsic features of data
  - RBMs “pre-train” MLP to avoid local minima!
  - RBM trained on unlabelled data!



# Deep AutoEncoders

---

- Consider a 4:2:4 MLP (ENCODER):



- Inputs:  $(0\ 0\ 0\ 1)$ ,  $(0\ 0\ 1\ 0)$ ,  $(0\ 1\ 0\ 0)$ ,  $(1\ 0\ 0\ 0)$
- Outputs:  $(0\ 0\ 0\ 1)$ ,  $(0\ 0\ 1\ 0)$ ,  $(0\ 1\ 0\ 0)$ ,  $(1\ 0\ 0\ 0)$

## Encoders

- 
- Encoder network "learns" the concept of binary representation of integers!
  - General case:  $2^n:n:2^n$  networks
  - Somehow the network is able to "extract" the most concise representation of the input!
  - The same trick works for "real numbers": PCA networks
  - It can be used for dimensionality reduction, data compression, and visualization.


## Training Encoders: variants

---

- Backpropagation (good for shallow architectures)
- Enforced "symmetry of weights"  
(upper part = mirror of lower part)
- **Denoising Autoencoders**: the input layers gets a "noisy version of true  $x$ "; the target contains the original  $x$  ("noise": e.g., 10% of randomly selected pixels set to 0)
- **Deep Autoencoders**: the first "half" is trained unsupervised as a DBN (a stack of RBMs); the upper part is the "mirror" of the lower part; final phase trained by backpropagation

# Deep Autoencoders

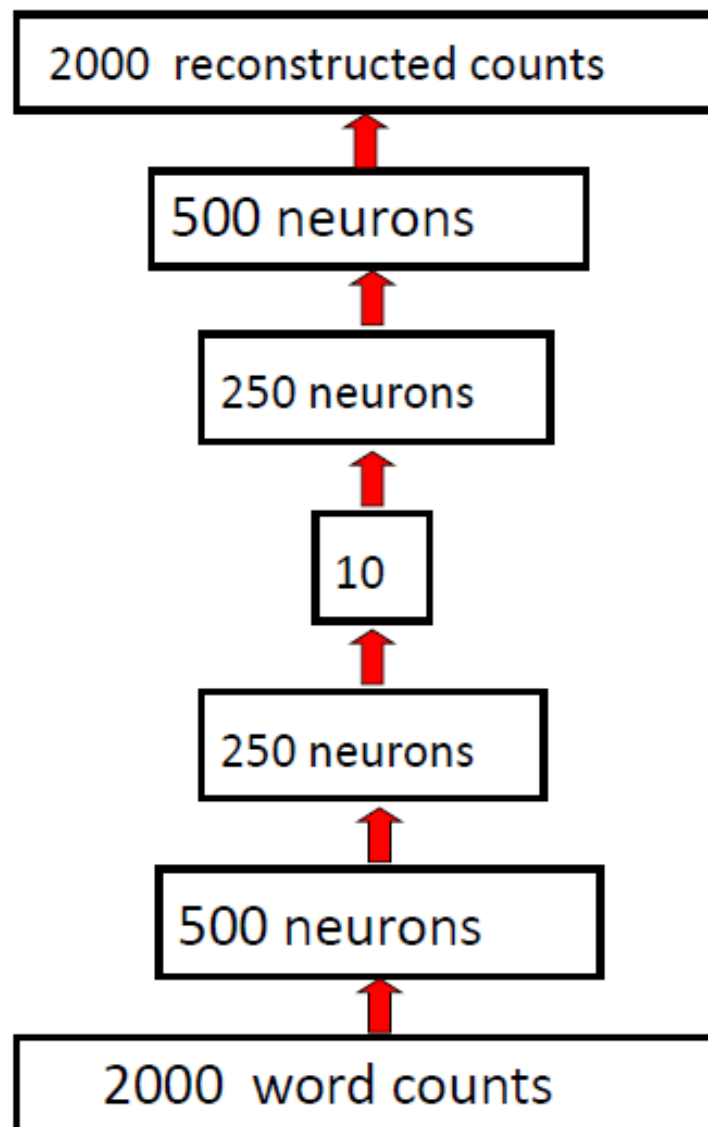
- They always looked like a really nice way to do non-linear dimensionality reduction:
  - But it is **very** difficult to optimize deep autoencoders using backpropagation.
- We now have a much better way to optimize them:
  - First train a stack of 4 RBM's
  - Then “unroll” them.
  - Then fine-tune with backprop.



## Applications: Classifying text documents

- A document can be characterized by the frequency of words that appear (ie, word counts for some dictionary become feature vector)
- Goals...
  1. Group/cluster similar documents
  2. Find similar documents

# How to compress the count vector

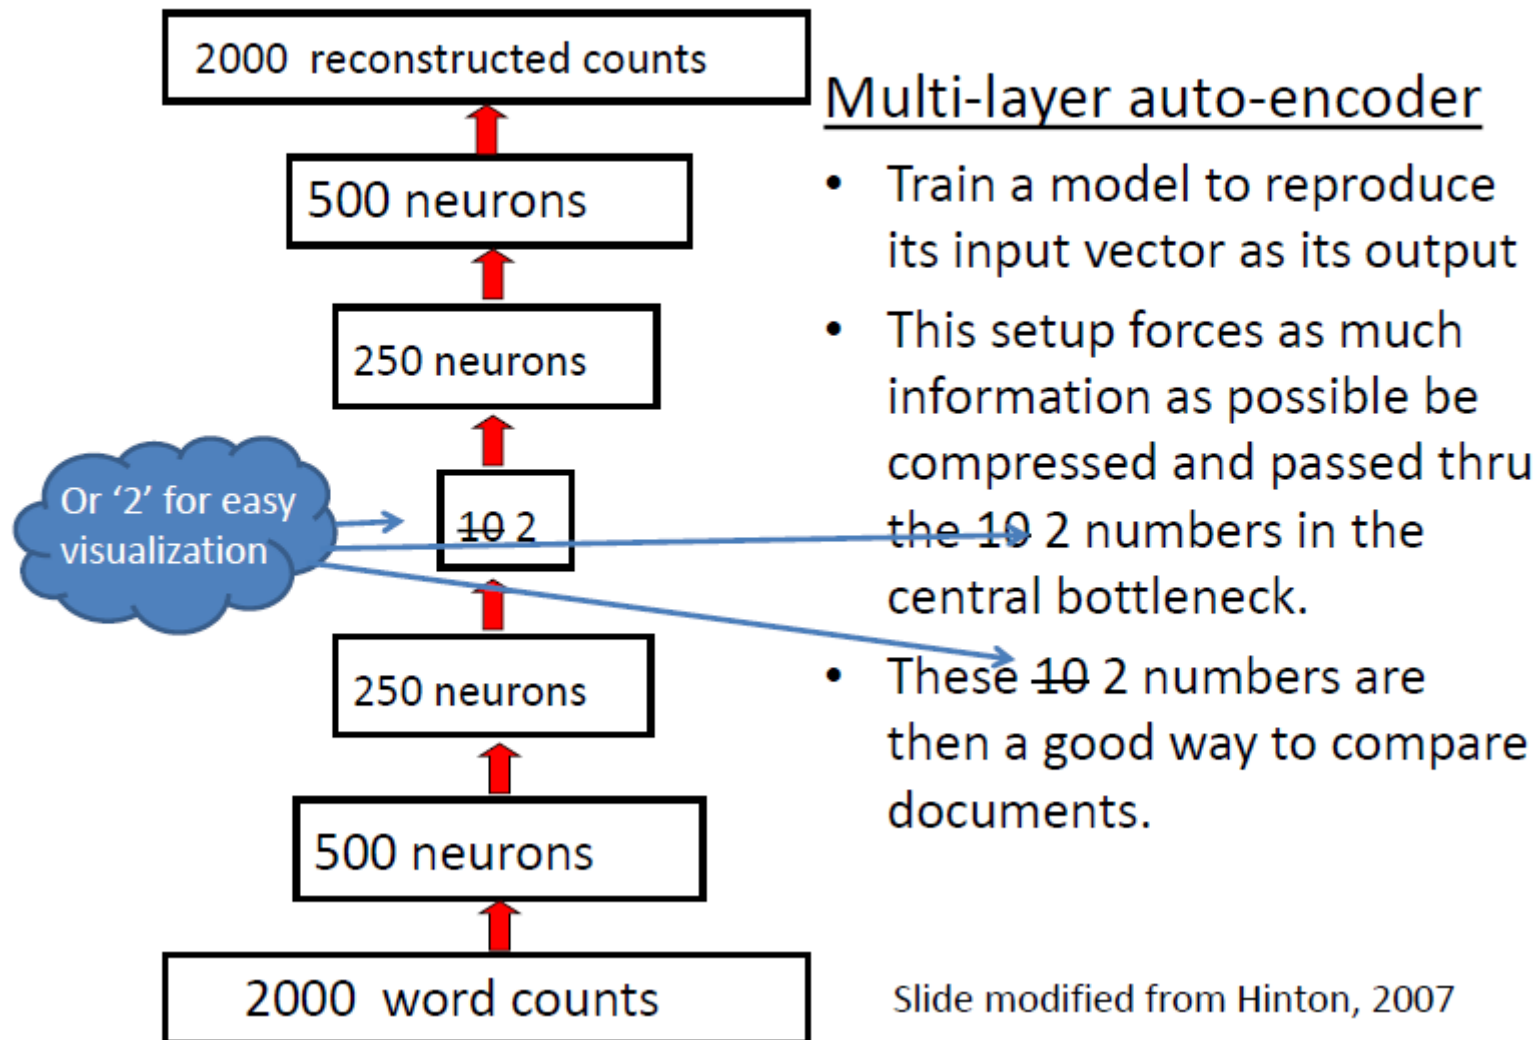


## Multi-layer auto-encoder

- Train a model to reproduce its input vector as its output
- This setup forces as much information as possible be compressed and passed thru the 10 numbers in the central bottleneck.
- These 10 numbers are then a good way to compare documents.


Slide modified from Hinton, 2007

# How to compress the count vector




# Clusters

LSA 2-D Topic Space




Autoencoder 2-D Topic Space



Images from Hinton, 2007



**Fig. 3.** (A) The two-dimensional codes for 500 digits of each class produced by taking the first two principal components of all 60,000 training images. (B) The two-dimensional codes found by a 784-1000-500-250-2 autoencoder. For an alternative visualization, see (8).



The same trick applied to "digits"

# RBM networks and Netflix Challenge

- **100.000.000** rating records collected over 1997-2005
- rating record:  
    <customer\_id, movie\_id, rating>
- **500.000** customers
- **18.000** movies
- rating = an integer: 1, 2, 3, 4 or 5
- Additionally, **3.000.000** test records:  
    <customer\_id, movie\_id, ? >

**GOAL: fill in “?’s” with numbers,  
so the error is minimized!**

**\$1.000.000 prize for improving Netflix system by 10%**

---

# Restricted Boltzmann Machines for Collaborative Filtering

---


Ruslan Salakhutdinov  
Andriy Mnih  
Geoffrey Hinton

University of Toronto, 6 King's College Rd., Toronto, Ontario M5S 3G4, Canada

RSALAKHU@CS.TORONTO.EDU

AMNIH@CS.TORONTO.EDU

HINTON@CS.TORONTO.EDU



output layer:  
e.g., 100 nodes

input layer:  
for each movie 5 nodes

# Google “zero-shot-translation”

<https://research.googleblog.com/2016/11/zero-shot-translation-with-googles.html>

